

# Artificial Intelligence

## DSE 3252

# Reinforcement Learning

---

ROHINI R RAO & PADMASHREE G

DEPT OF DATA SCIENCE & COMPUTER APPLICATIONS

JANUARY 2025

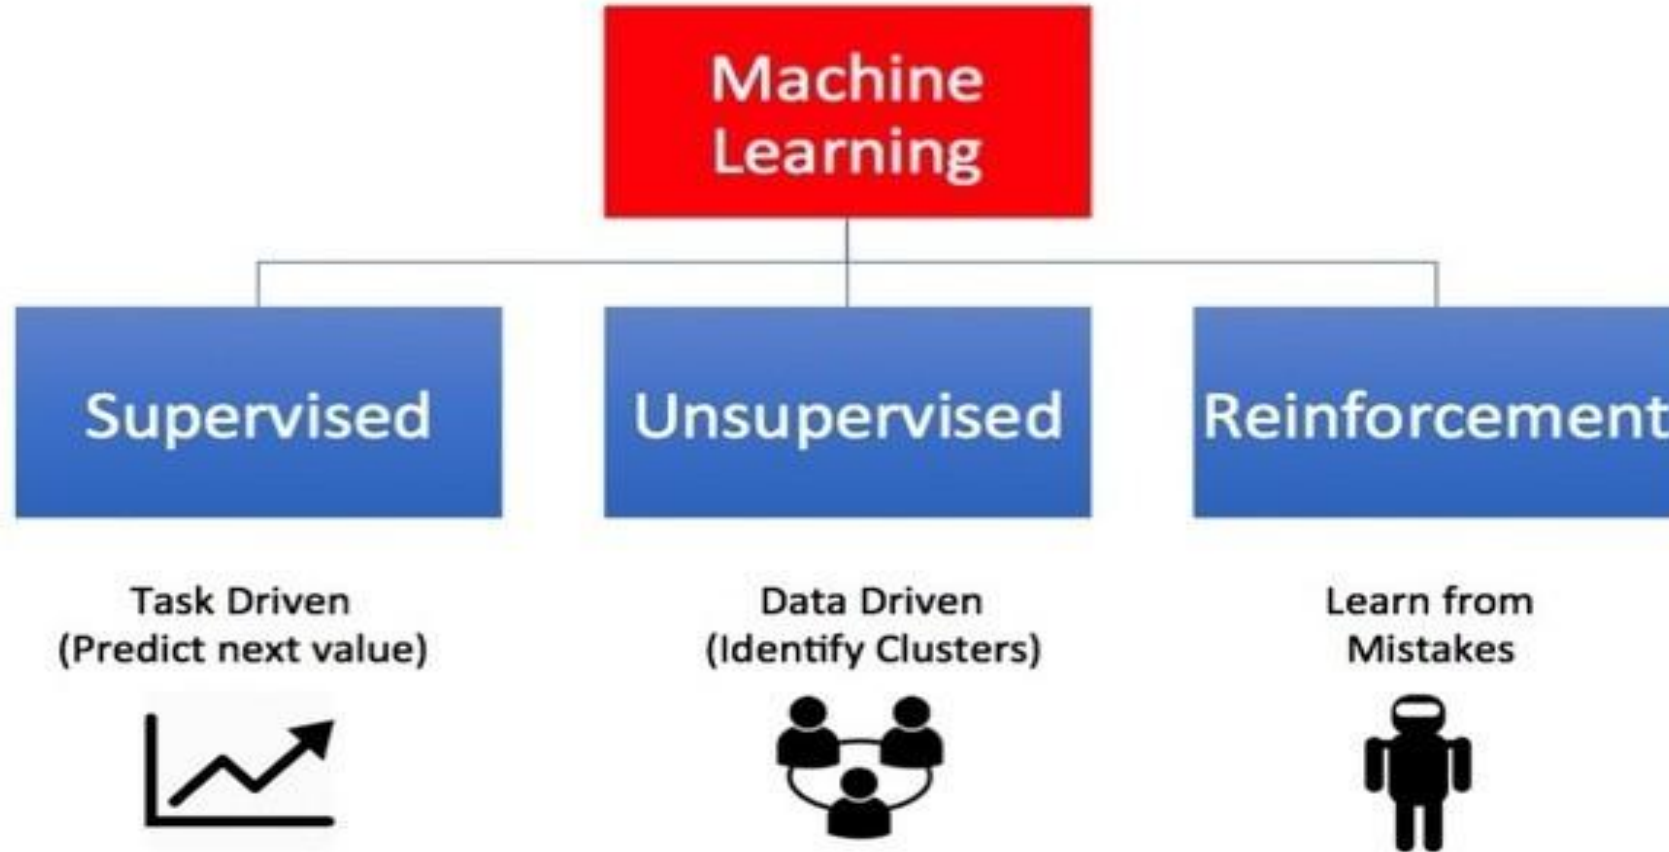


# What is Reinforcement Learning

---

- Reinforcement Learning is a
  - feedback-based Machine learning technique in which an agent learns to behave in an environment by performing the actions and seeing the results of actions.
  - For each good action, the agent gets positive feedback
  - for each bad action, the agent gets negative feedback or penalty.
- ***"Reinforcement learning is a type of machine learning method where an intelligent agent (computer program) interacts with the environment and learns to act within that."***


# Types of Machine Learning



In Reinforcement learning the goal is to find a suitable action model that would maximize the **total cumulative reward** of the agent

# Reinforcement Learning

---



# Problem Formulation in RL

## Environment

Physical world in which the agent operates

## State

- Current situation of the agent

## Reward


- Feedback from the environment

## Policy

- Method to map agent's state to actions

## Value

- Future reward that an agent would receive by taking an action in a particular state



# Markov Decision Process:

The mathematical framework for defining a solution in a reinforcement learning scenario

Can be designed as:

---

- **Set of states,  $S$**
- **Set of actions,  $A$**
- **Reward function,  $R$**
- **Policy,  $\pi$**
- **Value,  $V$**
- Set of action ( $A$ ) has to be taken to transition from our start state to our end state ( $S$ ).
- Model gets rewards ( $R$ ) (Positive or Negative) for each action we take
- The set of actions we took defines our policy ( $\pi$ )
- rewards we get in return define our value ( $V$ )
- Our task
  - *is to maximize our rewards by choosing the correct policy.*
  - *we have to maximize for all possible values of  $S$  for a time  $t$ .*

# Multi-Armed Bandit Problem

---

A bandit is defined as someone who steals your money.

One-armed bandit is a simple slot machine wherein you insert a coin into the machine, pull a lever, and get an immediate reward.

## **A multi-armed bandit**

- there are several levers that a gambler can pull, with each lever giving a different return.
- Probability distribution for the reward corresponding to each lever is different and is unknown to the gambler.
- The task is to identify which lever to pull to get maximum reward after a given set of trials.

# Multi-Armed Bandit problem (MAB)

---

- is a special case of **Reinforcement Learning**
- **MAB**
  - collects rewards in an environment by taking some actions after observing some state of the environment.
  - action taken by MAB does not influence the next state of the environment.
  - Therefore, MAB do not model state transitions, credit rewards to past actions, or "plan ahead" to get to reward-rich states.
- Goal of a MAB *agent* is to find a *policy* that collects as much reward as possible.
- **Exploration vs Exploitation Dilemma**
  - Not a good idea to exploit the action that promises the highest reward
  - because then there is a chance that we miss out on better actions if we do not explore enough.



# Value of Action

In our k-armed bandit problem, each of the k actions has an expected or mean reward given that that action is selected

---

- $A_t$  - action selected on time step t
- $R_t$  - corresponding reward as  $R_t$ .
- The value of an arbitrary action a, denoted  $q_*(a)$
- the expected reward given that a is selected:  $q_*(a) = E[R_t | A_t = a]$

If you knew the value of each action, then it would be trivial to solve the k-armed bandit

## Problem

- $Q_t(a)$  - the estimated value of action a at time step t
- $Q_t(a)$  to be close to  $q_*(a)$

## Exploitation

- maximize the expected reward on the one step
- maintain estimates of the action values
- Greedy action - at any time step there is at least one action whose estimated value is greatest

## Exploration

- exploration may produce the greater total reward in the long run
- greedy action's value is known with certainty, while several other actions are estimated to be nearly as good but with substantial uncertainty

# Action Value Methods

---

## Sample Average Method

$$Q_t(a) \doteq \frac{\text{sum of rewards when } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t} = \frac{\sum_{i=1}^{t-1} R_i \cdot \mathbb{1}_{A_i=a}}{\sum_{i=1}^{t-1} \mathbb{1}_{A_i=a}},$$

where  $\mathbb{1}_{\text{predicate}}$  denotes the random variable that is 1 if predicate is true and 0 if it is not

If the denominator

- is 0 - define  $Q_t(a)$  as some default value, such as 0.
- goes to infinity- by the law of large number,  $Q_t(a)$  converges to  $q_*(a)$

Greedy action selection method

$$A_t \doteq \operatorname{argmax}_a Q_t(a)$$

- $\operatorname{argmax}_a$  denotes the action  $a$  for which the expression that follows is maximized
- Variation -  $\epsilon$ -greedy - select randomly from among all the actions with equal probability

# exploration vs exploitation dilemma

## Pure exploitation approach

- select only one slot machine and keep pulling the lever all day long.
- may give you “some” payouts.
- might hit the jackpot (with a probability close to 0.00000....1)

## Pure exploration approach

- pull a lever of each & every slot machine
- Get sub-optimal payouts
- May be at least one of them would hit the jackpot.



## Exploration vs Exploitation trade-off

- To build an optimal policy, the agent faces the dilemma of exploring new states while maximizing its overall reward at the same time.

Arm	Reward
1	0
2	0
3	1
4	1
5	0
3	1
3	1
2	0
1	1
4	0
2	0

*The best overall strategy may involve short-term sacrifices.*

*Therefore, the agent should collect enough information to make the best overall decision in the future.*



# Exploitation vs Exploration

# $\epsilon$ -Greedy Method

---

- Behave greedily most of the time
- A few times with small probability Epsilon, select randomly from among all the actions with equal probability
- independently of the action-value estimates.

## **Advantages**

- as the number of steps increases, every action will be sampled an infinite number of times thus ensuring that all the  $Q_t(a)$  converge to  $q_*(a)$

# Incremental Implementation

---

$$\begin{aligned}Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\&= \frac{1}{n} \left( R_n + \sum_{i=1}^{n-1} R_i \right) \\&= \frac{1}{n} \left( R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i \right) \\&= \frac{1}{n} (R_n + (n-1)Q_n) \\&= \frac{1}{n} (R_n + nQ_n - Q_n) \\&= Q_n + \frac{1}{n} [R_n - Q_n],\end{aligned}$$

- Error in the estimate – [Target–OldEstimate]
- Target is the  $n^{\text{th}}$  reward.
- StepSize changes from time step to time step
- In processing the  $n^{\text{th}}$  reward for action  $a$ , the method uses the step-size parameter  $1/n$
- step size parameter is denoted as  $\alpha_t(a)$
- Update Rule

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} \left[ \text{Target} - \text{OldEstimate} \right]$$

# Simple Bandits Algorithm

---

Initialize, for  $a = 1$  to  $k$ :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Loop forever:


$$A \leftarrow \begin{cases} \operatorname{argmax}_a Q(a) & \text{with probability } 1 - \varepsilon \\ \text{a random action} & \text{with probability } \varepsilon \end{cases} \quad (\text{breaking ties randomly})$$

$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$

$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

# Average Performance of $\epsilon$ -greedy





# Strategies to encourage exploration

## - Optimistic initial action values

---

There are 3 arms with unknown distributions


$R_1 \sim N(1,1)$

$R_2 \sim N(2,1)$  (optimal arm)

$R_3 \sim N(1.5,1)$

Step (t)	Selected Arm (a)	Observed Reward (R)	Update $Q_t(a)$	New Q values
0			$Q_1=Q_2=Q_3=5$	( 5, 5, 5)
1	1 (random)	0.5	$Q_1=5+(1/1)(0.5-5)$	(0.5,5, 5)
2	2 (greedy)	2	$Q_2= 5 + (1/1)(2-5)$	(0.5,2, 5)
3	3 (greedy)	1.3	$Q_3 = ?$	(0.5,2,
4	..			
5	...			

# The effect of optimistic initial action values




# Upper-confidence-bound action selection

## Optimism in the face of uncertainty

$$A_t \doteq \operatorname{argmax}_a \left[ Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right],$$

where;

- $Q_t(a)$  is the estimated value of action 'a' at time step 't'.
- $N_t(a)$  is the number of times that action 'a' has been selected, prior to time 't'.
- 'c' is a confidence value that controls the level of exploration.



# UCB Action Selection

$$A_t \doteq \operatorname{argmax}_a \left[ Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right],$$

## **Exploitation:**

- $Q_t(a)$  represents the exploitation
- if you don't know which action is best then choose the one that currently looks to be the best

## **Exploration:**

- If an action hasn't been tried very often, or not at all, then  $N_t(a)$  will be small. Consequently, the uncertainty term will be large, making this action to be selected
- As  $N_t(a)$  increments, and the uncertainty term decreases, making it less likely that this action will be selected as a result of exploration
- it may still be selected as the action with the highest value

***As  $n$  goes to infinity the exploration term gradually decreases until eventually, actions are selected based only on the exploitation term.***

# Steps followed in UCB agent

---

*1. At each round  $t$ , we compute two numbers for arm  $A$ .*

->  $N_A(t)$  = number of times the arm  $A$  was selected up to round  $t$ .

->  $R_A(t)$  = number of rewards of the arm  $A$  up to round  $t$ .

*2. From these two numbers we have to calculate,*

a. The average reward of machine  $m$  up to round  $t$

$$r_A(t) = R_A(t) / N_A(t).$$


b. The confidence interval  $[ r_A(t) - \Delta_A(t), r_A(t) + \Delta_A(t) ]$  at round  $n$

with,  $\Delta_A(t) \Rightarrow c * \text{sqrt}( \ln(t) / N_a(t) )$

*1. We select the arm  $A$  that has the maximum UCB,  $( r_A(t) + \Delta_A(t) )$*

# Upper-confidence-bound action selection

---



# Gradient Bandit Algorithms

---

Probability of taking action  $a$  at time  $t$

$$\Pr\{A_t = a\} \doteq \frac{e^{H_t(a)}}{\sum_{b=1}^k e^{H_t(b)}} \doteq \pi_t(a)$$

The action preferences are updated by

$$\begin{aligned} H_{t+1}(A_t) &\doteq H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)), & \text{and} \\ H_{t+1}(a) &\doteq H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a), & \text{for all } a \neq A_t \end{aligned}$$

Initially all action preferences are the same (e.g.,  $H_1(a) = 0$ , for all  $a$ )

As Stochastic Gradient Ascent

$$H_{t+1}(a) \doteq H_t(a) + \alpha \frac{\partial \mathbb{E}[R_t]}{\partial H_t(a)}$$

Where


$\alpha > 0$  is step size parameter

$$\bar{R}_t \in \mathbb{R}$$

Is average of all rewards upto and including time  $t$

$$\mathbb{E}[R_t] = \sum_x \pi_t(x) q_*(x)$$

# Gradient Bandit Algorithms



Without Baseline –

$\bar{R}_t$  is set to 0



# Ad Optimization

---

Ad 1	Ad 2	Ad 3	Ad 4	Ad 5	Ad 6	Ad 7	Ad 8	Ad 9	Ad 10
1	0	0	0	1	0	0	0	1	0
0	0	0	0	0	0	0	0	1	0
0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
1	1	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	1	0

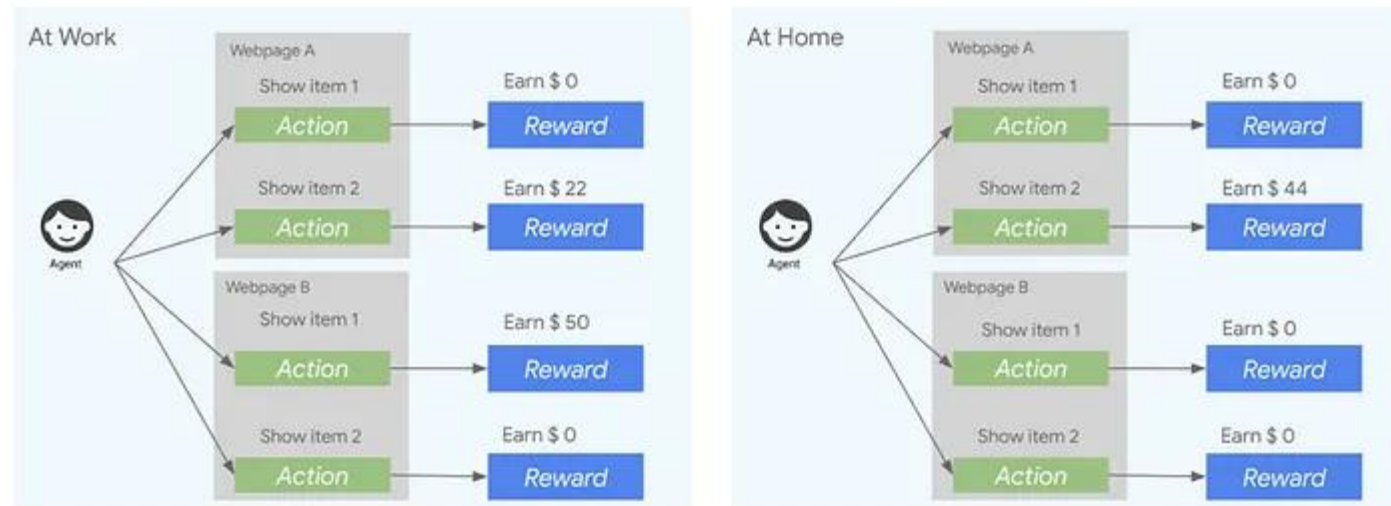
# Associative Search (Contextual Bandits)

In k-armed bandit problem each action affects only the immediate reward

## Contextual Bandits

- actions are affected by **context**, which in turn affects the reward
- In a general RL task, there is more than one situation/context
- **the goal is to learn a policy**: a mapping from situations/context to the actions that are best in those situations

### Contextual bandits



# Multi-armed Bandits Problem

---

- K actions (feature-free)
- Each action has an average reward (unknown):  $\mu_k$
- For  $t=1, \dots, T$  (unknown)
  - Choose an action  $a_t$  from  $\{1, \dots, K\}$  actions
  - Observe a random reward  $y_t$ , where  $y_t$  is bounded  $[0,1]$
  - $E[y_t] = \mu_{a,t}$  : Expected reward of action  $a_t$
- Minimizing Regret: 
$$R = \sum_{t=1}^T [\mu^* - \mu_{a,t}]$$

**regret** is the difference between the total reward achieved by always selecting the optimal arm and the total reward achieved by the algorithm.

**Q. How to choose an action to minimize regret?**

# Feature free bandit setting

---

Users  $u_1$  with age YOUNG  
and  $u_2$  with age OLD



$u_1$



$u_2$

**Retirement planning wishes  
vs. reality**

**The Player** Wizarding World of Harry Potter  
ride may conjure a new path for theme park  
rides

**Elon Musk: 198,000 Tesla Model 3 Orders  
Received in 24 Hours**

**Not tired yet: Warriors top Spurs for  
72nd win, set up date with history**

# Contextual Bandit Problem

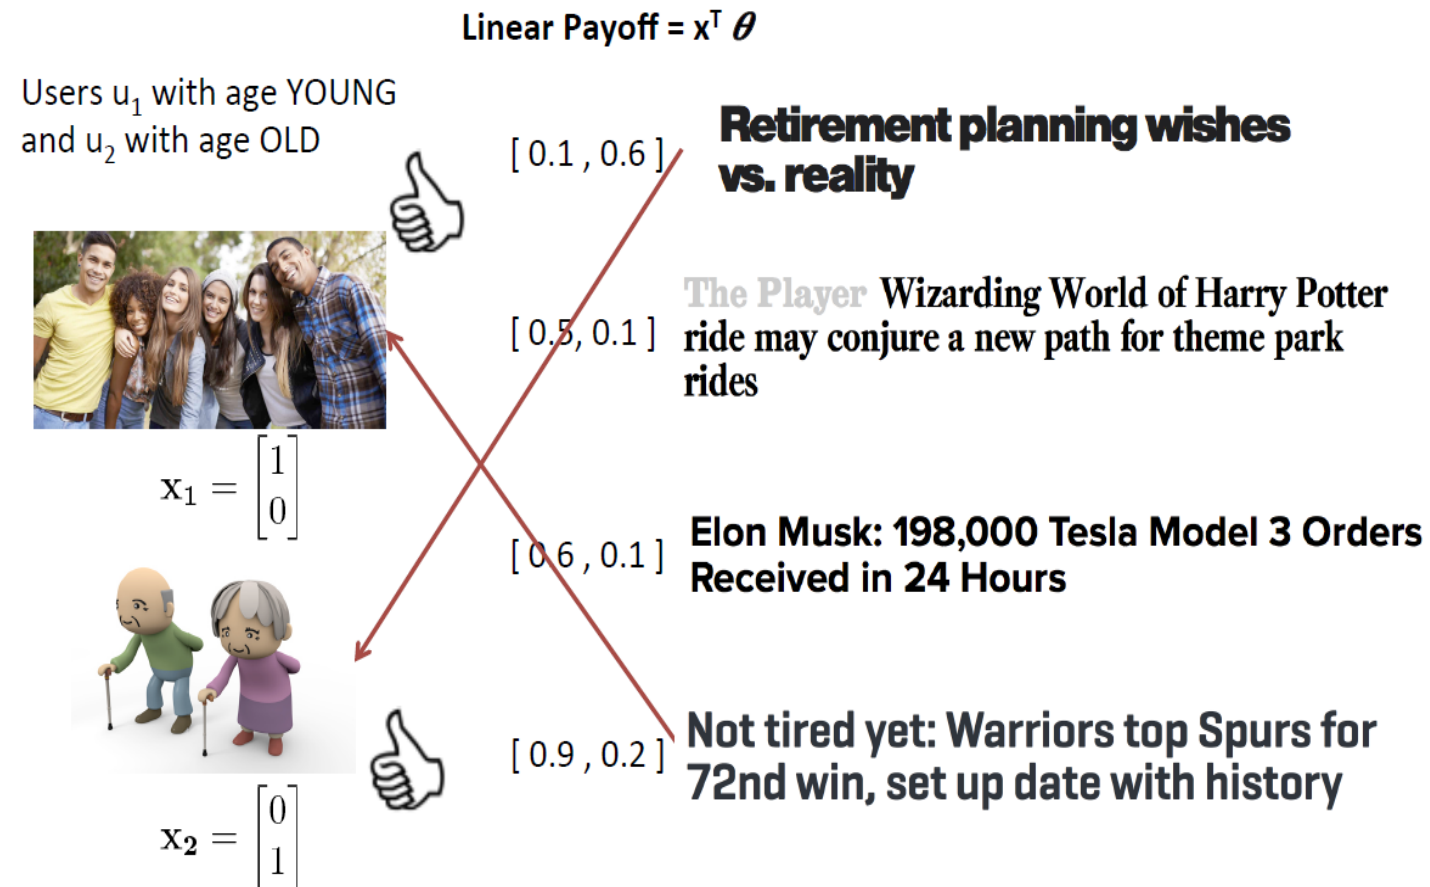
- For  $t=1,\dots,T$  (unknown)
  - User  $u_t$ , set  $A_t$  of actions (a)
  - Feature vector (context)  $\mathbf{x}_{t,a}$  : summarizes both user  $u_t$  and action a
  - Based on previous results, choose  $a_t$  from  $A_t$
  - Receive payoff  $r_{t,a_t}$
  - Improve selection strategy with new observation set  $(x_{t,a_t}, a_t, r_{t,a_t})$

$$\mathbf{E}[r_{t,a} | \mathbf{x}_{t,a}] = \mathbf{x}_{t,a}^\top \boldsymbol{\theta}_a^*.$$

Minimizing Regret:  $R(T) = \mathbf{E} \left[ \sum_{t=1}^T \left( r_{t,a_t^*} - r_{t,a_t} \right) \right]$

Action with maximum  
expected payoff at time t

# Contextual Linear Bandits setting



## Contextual Bandit

For each trial  $t=1,2,3,\dots, T$

1. Observe environment  $x_{t,a} \in \mathbb{R}^d$ , i.e. user  $u_t$  a set of actions  $\mathcal{A}_t$  and both their features
2. Choose an arm  $a_t \in \mathcal{A}$  based on previous trials and receive payoff  $r_{t,a_t}$ .
3. Improve arm-selection strategy with new observation  $(x_{t,a_t}, a_t, r_{t,a_t})$



Minimize expected regret, i.e

$$R_A(T) = \mathbb{E} \left[ \sum_{t=1}^T r_{t,a_t^*} \right] - \mathbb{E} \left[ \sum_{t=1}^T r_{t,a_t} \right]$$

## Example: News Recommendation

For each time the news page is loaded  $t=1,2,3,\dots, T$

1. Arms or actions are the articles, which can be shown to the user. The environment could be user and article information.
2. If the article is clicked  $r_{t,a_t} = 1$  otherwise 0.
3. Improve new article selection





# Linear Disjoint Model

- *disjoint* since the parameters are not shared among different arms.
- To solve for the coefficient vector  $\theta$  ridge regression is applied to the training data.

$$E[r_{t,a} | x_{t,a}] = [x_{t,a}]^T \theta_a^*$$

▪ How to estimate  $\theta_a$ ?

▪ Linear regression solution to  $\theta_a$  is

$$\widehat{\theta}_a = \underset{\theta}{\operatorname{argmin}} \sum_{m \in D_a} ([x_{t,a}]^T \theta_a - b_a^{(m)})^2$$

We can get:

$$\widehat{\theta}_a = (D_a^T D_a + I_d)^{-1} D_a^T b_a$$

$D_a$  is a  $m \times d$  matrix of  $m$  training inputs  $[x_{t,a}]$

$b_a$  is a  $m$ -dimension vector of responses to  $a$  (click/no-click)



# linUCB Algorithm

- Initialization:

$$A_a \stackrel{\text{def}}{=} \mathbf{D}_a^T \mathbf{D}_a + \mathbf{I}_d$$

- For each arm  $a$ :

- $A_a = \mathbf{I}_d$

//identity matrix  $d \times d$

- $b_a = [0]_d$

//vector of zeros

- Online algorithm:

- For  $t=[1:T]$ :

- Observe features for all arms  $a : x_{t,a} \in R^d$

- For each arm  $a$  :

- $\theta_a = A_a^{-1} b_a$

//regression coefficients

- $p_{t,a} = [x_{t,a}]^T \theta_a + \alpha \sqrt{[x_{t,a}]^T A_a^{-1} x_{t,a}}$

- Choose arm  $a_t = \operatorname{argmax}_a p_{t,a}$

//choose arm

- $A_{a_t} = A_{a_t} + x_{t,a_t} [x_{t,a_t}]^T$

//update A for the chosen arm  $a_t$

- $b_{a_t} = b_{a_t} + r_t x_{t,a_t}$

//update b for the chosen arm  $a_t$

# Thomson Sampling

---

## Basic Intuition

1. all machines are assumed to have a uniform distribution of the probability of reward
  2. For each observation, a new distribution of rewards is generated (exploration)
  3. Further observations are used to update the success distributions of rewards
  4. After sufficient observations, each slot machine will have a success distribution of rewards (exploitation)
- A simple natural Bayesian heuristic
    - Maintain a belief(distribution) for the unknown parameters
    - Each time, pull arm  $a$  and observe a reward  $r$
  - Initialize priors using belief distribution
    - For  $t=1:T$ :
      - Sample random variable  $X$  from each arm's belief distribution
      - Select the arm with largest  $X$
      - Observe the result of selected arm
      - Update prior belief distribution for selected arm

# Example: Web Content Personalization

---

## Vowpal Wabbit

- an interactive ML library and the RL framework for services like Microsoft Personalizer.
- It allows for maximum throughput and lowest latency when making personalization ranks and training the model with all events

**Con-Ban Agent** performs the following functions:

Some context 'x' arrives and is observed by Con-Ban Agent.

Con-Ban Agent chooses an action 'a' from a set of actions A, i.e.,  $a \in A$  (A may depend on 'x').

Some reward 'r' for the chosen 'a' is observed by Con-Ban Agent.

For example: **Con-Ban Agent news website:**

- **Decision to optimize:** articles to display to user.
- **Context:** user data (browsing history, location, device, time of day)
- **Actions:** available news articles
- **Reward:** user engagement (click or no click)

# Summary of Notation

$\doteq$	equality relationship that is true by definition
$\approx$	approximately equal
$\propto$	proportional to
$\Pr\{X=x\}$	probability that a random variable $X$ takes on the value $x$
$X \sim p$	random variable $X$ selected from distribution $p(x) \doteq \Pr\{X=x\}$
$\mathbb{E}[X]$	expectation of a random variable $X$ , i.e., $\mathbb{E}[X] \doteq \sum_x p(x)x$
$\operatorname{argmax}_a f(a)$	a value of $a$ at which $f(a)$ takes its maximal value
$\ln x$	natural logarithm of $x$
$e^x$	the base of the natural logarithm, $e \approx 2.71828$ , carried to power $x$ ; $e^{\ln x} = x$
$\mathbb{R}$	set of real numbers
$f: \mathcal{X} \rightarrow \mathcal{Y}$	function $f$ from elements of set $\mathcal{X}$ to elements of set $\mathcal{Y}$
$\leftarrow$	assignment
$(a, b]$	the real interval between $a$ and $b$ including $b$ but not including $a$
$\varepsilon$	probability of taking a random action in an $\varepsilon$ -greedy policy
$\alpha, \beta$	step-size parameters
$\gamma$	discount-rate parameter
$\lambda$	decay-rate parameter for eligibility traces
$\mathbb{1}_{\text{predicate}}$	indicator function ( $\mathbb{1}_{\text{predicate}} \doteq 1$ if the <i>predicate</i> is true, else 0)

In a multi-arm bandit problem:

$k$	number of actions (arms)
$t$	discrete time step or play number
$q_*(a)$	true value (expected reward) of action $a$
$Q_t(a)$	estimate at time $t$ of $q_*(a)$
$N_t(a)$	number of times action $a$ has been selected up prior to time $t$
$H_t(a)$	learned preference for selecting action $a$ at time $t$
$\pi_t(a)$	probability of selecting action $a$ at time $t$
$\bar{R}_t$	estimate at time $t$ of the expected reward given $\pi_t$

# References

---

1. Richard S Sutton, Andrew G Barto, Reinforcement Learning, second edition, MIT Press
2. Russell S., and Norvig P., Artificial Intelligence A Modern Approach (3e), Pearson 2010
3. <https://www.coursera.org/learn/fundamentals-of-reinforcement-learning/home/week/1>